



UNIVERSITY OF ASIA PACIFIC (UAP)
Department of Computer Science and Engineering

A BUSINESS CONCEPT REPORT ON

KARIGOR AI

AI-Powered Quality Inspection for Small Manufacturers

Developed in Response to the Bangladesh AI Business Summit 2026 Student EoI Call

PREPARED BY

Noshin Tabassum

Student ID: 24101035 | Section: A | Semester: 3.1

Department of Computer Science and Engineering
University of Asia Pacific (UAP)
Dhaka, Bangladesh

August 2026

Executive Summary

Product quality is one of the most decisive factors in the competitiveness of any manufacturing business. For small and medium-sized manufacturers, however, maintaining consistent quality control is often difficult. Most small factories and workshops rely on manual visual inspection, in which human inspectors examine finished or semi-finished products for visible defects. While manual inspection is flexible and low in initial cost, it is inherently limited by human fatigue, subjective judgment, and the difficulty of maintaining digital, analyzable quality records over time. As production volumes grow, these limitations become more pronounced, and inconsistent inspection can lead to defective products reaching customers or to unnecessary rework and waste.

At the same time, industrial machine-vision systems capable of automated visual inspection exist, but these systems are typically expensive, require specialized technical personnel, and are designed around the scale and integration requirements of large factories. This creates a gap for small and medium-sized manufacturers, who face real quality-control challenges but lack access to inspection technology suited to their scale, budget, and technical capacity.

Karigor AI is proposed as a conceptual response to this gap. It envisions an affordable, modular, AI-assisted visual quality-inspection solution built specifically around the constraints of small manufacturers. The proposed system would use a camera positioned at a product inspection point, together with a computer-vision model, to examine manufactured items for visible defects. Depending on the product and defect type, the system could provide a pass/fail decision, classify the defect where sufficient training data exist, and optionally localize the defect on the product image. Every prediction would carry a confidence score, and predictions below a defined confidence threshold would be routed to a human inspector for final judgment. This human-in-the-loop design is a deliberate and central part of the concept: Karigor AI is intended to assist and augment human quality-control personnel, not to replace them, particularly given the serious consequences of an AI system incorrectly accepting a defective product.

Beyond defect detection itself, Karigor AI envisions a simple dashboard that would allow factory owners and quality managers to view inspection history, batch-level summaries, and basic defect-trend analytics — capabilities that are difficult to achieve consistently with paper-based or purely manual inspection records. The target users are small and medium-sized manufacturers across sectors where products can be evaluated visually, such as garments and textile accessories, packaging, plastic goods, leather products, handicrafts, and light manufacturing components. Because a single computer-vision model cannot reasonably be expected to inspect every kind of product, the concept assumes that each deployment would require product-specific image data, calibration, and validation.

From a business perspective, Karigor AI is conceived around a modular adoption path: an initial setup and calibration fee, followed by a recurring subscription covering software, dashboard access, and model maintenance, with optional future revenue from custom model training, additional production-line coverage, and technical support. This structure is intended to lower the upfront barrier to adoption relative to full industrial automation, allowing a manufacturer to begin with a single inspection point before expanding.

It is important to state clearly that Karigor AI, as presented in this report, is a conceptual business and technical proposal. It has not been implemented as a working prototype, has not been tested on real production data, and no performance results, customer relationships, or commercial outcomes currently exist. The report describes a proposed architecture, a plausible development methodology grounded in established computer-vision research, and a realistic pathway toward a future minimum viable product and pilot validation with a cooperating small manufacturer. The

intention of this document is to demonstrate a technically grounded and commercially plausible direction for affordable AI-assisted quality inspection, while being transparent about what remains to be built, tested, and validated before Karigor AI could operate as a real product.

Project Context Note

Karigor AI was independently developed in response to the student Expression of Interest call associated with the Bangladesh AI Business Summit 2026, organized by the Dhaka Chamber of Commerce & Industry (DCCI) [9]. The concept explores the potential use of AI-powered visual inspection to support quality-control activities in small manufacturing environments. The concept remained at the independent conceptual-development stage and was not formally submitted, selected, presented, or exhibited at the event. Accordingly, this report presents the proposed technical architecture, business model, implementation pathway, and future validation strategy as an independent business and technology concept rather than as a completed commercial system.

Table of Contents

Executive Summary.....	2
Project Context Note.....	3
List of Figures.....	4
List of Tables.....	4
Chapter 1: Introduction.....	5
Chapter 2: Problem Statement, Aim and Objectives.....	6
Chapter 3: Proposed Karigor AI Solution.....	8
Chapter 4: Target Users, Use Cases and Value Proposition.....	10
Chapter 5: Proposed Technical Architecture and AI Methodology.....	12
Chapter 6: Business Model and Market Opportunity.....	15
Chapter 7: Competitive Positioning and Feasibility.....	18
Chapter 8: SWOT, Risks and Responsible AI.....	19
Chapter 9: Sustainability and Potential Impact.....	21
Chapter 10: Implementation Roadmap and MVP.....	22
Chapter 11: Limitations and Future Development.....	24
Chapter 12: Conclusion.....	25
References.....	26

List of Figures

Figure 1. Proposed Operational Workflow of Karigor AI.....	10
Figure 2. Conceptual System Architecture of Karigor AI.....	13
Figure 3. Proposed Business Model Canvas for Karigor AI.....	17
Figure 4. Proposed Implementation Roadmap for Karigor AI.....	22

List of Tables

Table 1. Target Users and Use Cases.....	12
Table 2. Proposed Evaluation Metrics for a Future Karigor AI Prototype.....	14
Table 3. Proposed Business Revenue Channels.....	16
Table 4. Comparison with Existing Quality Inspection Approaches.....	18
Table 5. SWOT Analysis.....	19
Table 6. Risk and Mitigation Matrix.....	20
Table 7. Proposed Implementation Phases.....	22

Chapter 1: Introduction

1.1 Background

Consistent product quality is a fundamental requirement for any manufacturing business, directly affecting customer satisfaction, brand reputation, and long-term profitability. Manufacturers of all sizes rely on quality-control (QC) processes to identify defective products before they reach the market. Among the various forms of quality control, visual inspection remains one of the most widely used methods, since many manufacturing defects — such as surface scratches, stitching faults, discoloration, misalignment, or packaging irregularities — are directly observable by eye.

In most small and medium-sized manufacturing environments, visual inspection is performed manually by human inspectors. This approach has clear advantages: it requires no specialized equipment, it is flexible across product types, and human inspectors can apply contextual judgment that rule-based systems may lack. However, research in human-factors engineering has consistently shown that human visual inspection performance degrades with fatigue, repetitive workload, and inspection duration, and that inspection accuracy can vary considerably between individual inspectors and across working conditions [4]. These characteristics make manual inspection difficult to scale reliably as production volume grows.

Automated visual inspection using computer vision and deep learning has advanced substantially in recent years. Convolutional neural networks (CNNs) and object-detection architectures such as the YOLO family have demonstrated strong performance in detecting and classifying surface and structural defects across a range of industrial materials, including steel, textiles, and manufactured components [5], [6]. These advances have made automated visual inspection technically feasible even in settings with only moderate amounts of training data, particularly when combined with transfer learning from models pretrained on large image datasets.

1.2 Motivation

Despite these technical advances, automated visual inspection remains largely inaccessible to small and medium-sized manufacturers. Established industrial machine-vision systems are typically designed for large-scale, high-throughput factories, and are priced, integrated, and supported accordingly. Small manufacturers frequently lack the capital budget, in-house technical expertise, and production volume that would justify such systems. At the same time, research on Industry 4.0 adoption in Bangladesh and other developing economies consistently identifies high implementation cost, limited technical skills, and infrastructure constraints as the most significant barriers preventing SMEs from adopting advanced manufacturing technologies [1], [2], [3].

This gap between the technical feasibility of AI-assisted visual inspection and its practical inaccessibility to small manufacturers is the central motivation for Karigor AI. Small manufacturers in sectors such as garments, packaging, plastics, leather goods, and light consumer products face real quality-control challenges, yet they are largely excluded from the benefits of automated inspection technology that larger competitors can afford. An affordable, modular, and appropriately scoped AI-assisted inspection concept, designed specifically around the constraints of small manufacturers rather than adapted from large-factory systems, could help narrow this gap.

Bangladesh's manufacturing sector, and particularly its SME segment, illustrates this opportunity clearly. SMEs form the overwhelming majority of manufacturing establishments in Bangladesh and are a major source of industrial employment [7], [8]. Strengthening the quality-

control capacity of this segment, even modestly, could have a meaningful cumulative effect on product consistency, waste reduction, and competitiveness across the wider manufacturing economy — motivating the exploration of low-cost, locally applicable AI inspection concepts such as Karigor AI.

1.3 Proposed Concept

Karigor AI proposes an AI-powered visual quality-inspection concept intended specifically for small and medium-sized manufacturers. The proposed system centers on a camera positioned at a fixed inspection point, capturing images of products as they move through or pause at that point. These images would be processed by a computer-vision model trained to recognize visible defects relevant to the specific product category being inspected. Depending on the deployment and available training data, the system could provide a pass/fail decision, classify the type of defect detected, and, where appropriate, indicate the approximate location of the defect on the product image.

A defining characteristic of the Karigor AI concept is its human-in-the-loop design. Rather than presenting AI predictions as final and autonomous decisions, the system is intended to attach a confidence score to each prediction. Predictions with high confidence could proceed directly to a pass or flag decision, while predictions falling below a defined confidence threshold would be routed to a human inspector for verification. This design reflects a conscious choice to treat Karigor AI as a decision-support tool that augments human quality-control judgment, rather than an infallible autonomous inspector — a distinction that is particularly important given that an undetected defect (a false accept) can have more serious consequences than an unnecessary manual review.

The concept also includes a dashboard and record-keeping layer intended to give factory owners and quality managers visibility into inspection outcomes over time — including pass/fail rates, defect-type trends, and batch-level summaries — supporting more informed quality-management decisions than paper-based or purely manual record-keeping typically allows.

1.4 Scope of the Report

This report presents Karigor AI as a conceptual business and technology proposal. It describes the problem the concept addresses, a proposed technical architecture and AI development methodology grounded in established computer-vision literature, target users and illustrative use cases, a proposed business model and market context, competitive positioning, feasibility considerations, risk and ethical analysis, a proposed implementation roadmap toward a minimum viable product, and an honest account of the concept's current limitations.

Karigor AI has not been implemented as a working prototype and has not been tested on real manufacturing data. No performance results, pilot outcomes, customer relationships, revenue, or commercial partnerships currently exist. All technical and business content in this report should be read as a proposed direction supported by relevant research and industry context, rather than as a description of a deployed or validated system.

Chapter 2: Problem Statement, Aim and Objectives

2.1 Problem Statement

Small and medium-sized manufacturers face a persistent quality-control problem that stems from two intersecting difficulties: the inherent limitations of manual visual inspection, and the practical inaccessibility of automated alternatives.

Limitations of Manual Visual Inspection

- Human fatigue: Inspection accuracy declines measurably over the course of a shift, as sustained visual attention is cognitively demanding and prone to lapses [4].
- Subjectivity and inconsistency: Different inspectors, or the same inspector at different times, may apply inspection criteria inconsistently, leading to variable pass/fail outcomes for similar products [4].
- Limited throughput: Manual inspection speed is constrained by human processing capacity, which can create bottlenecks as production volume increases.
- Overlooked defects: Repetitive, low-variation tasks make it easy for inspectors to miss defects, particularly subtle or infrequent ones.
- Weak digital traceability: Manual inspection records are frequently kept on paper or not kept in a structured, analyzable form at all, making it difficult to identify recurring defect patterns over time.

Barriers Facing Small Manufacturers

- High cost of industrial machine-vision systems, which are typically priced and engineered for large-scale factory deployment [3].
- Lack of in-house technical personnel capable of installing, calibrating, or maintaining specialized inspection equipment.
- Integration complexity, since many commercial systems assume compatibility with existing factory automation infrastructure that small manufacturers may not have [2].
- Difficulty assembling large, labelled image datasets required to train conventional deep-learning models from scratch.
- A general need for simple deployment, local language support, and accessible technical assistance suited to smaller operations.

These two sets of difficulties reinforce each other: because automated alternatives are inaccessible, small manufacturers remain dependent on manual inspection, and because manual inspection is inherently limited, quality inconsistency, rework, and undetected defects remain persistent operational challenges. This creates a clear need for an inspection approach that is technically capable enough to meaningfully assist quality control, yet affordable, modular, and simple enough to be realistically adopted by a small manufacturer. It is important to note that the intention is not to eliminate human inspectors, but to support them — combining the consistency and speed of an AI-assisted first pass with the contextual judgment of a human reviewer for uncertain cases.

2.2 Project Aim

The aim of this project is: ***“To conceptualize an affordable AI-powered visual quality-inspection solution that can assist small manufacturers in detecting product defects, improving inspection consistency, and maintaining digital quality-control records.”***

2.3 Specific Objectives

- To design a conceptual AI-assisted defect-detection workflow that combines computer-vision analysis with human oversight for uncertain predictions.
- To propose a system architecture and hardware configuration that remains affordable and realistic for small manufacturers with limited technical capacity.
- To define a modular deployment approach that allows a manufacturer to begin with a single product category and inspection point before expanding.

- To ensure the proposed interface and workflow are simple enough for factory staff without specialized AI training to operate and interpret.
- To incorporate digital inspection records, reporting, and basic analytics that support traceability and identification of recurring defect patterns.
- To outline a realistic pathway toward a future minimum viable product and a small-scale pilot validation with a cooperating manufacturer.

Chapter 3: Proposed Karigor AI Solution

3.1 Overview

Karigor AI is conceptualized as a layered system consisting of a physical inspection setup, an AI processing pipeline, an application and dashboard layer, and a human-in-the-loop decision-support mechanism. Each layer is described below, followed by the proposed operational workflow shown in Figure 1.

3.2 Physical Inspection Layer

At the physical level, Karigor AI envisions a single, fixed inspection point at which products are presented to a camera, either manually by an operator or as part of an existing production step. The proposed setup consists of a standard industrial or USB camera, optional controlled lighting to reduce inconsistency caused by variable ambient light, a simple product placement fixture to keep imaging conditions consistent, and a low-cost computing device — such as a small PC or edge-computing unit — to run the inspection software. This configuration is intended to be inexpensive relative to dedicated industrial machine-vision hardware, and to require minimal changes to a manufacturer's existing production layout.

3.3 AI Processing Layer

The AI processing layer would receive the captured image and pass it through a preprocessing stage — resizing, cropping, and normalization — before feeding it into a trained computer-vision model. Depending on the product and defect type, the model could perform binary classification (acceptable vs. defective), multi-class defect classification, object detection to localize defects, or anomaly detection where labelled defect examples are scarce. Each prediction would be accompanied by a confidence score, which forms the basis for the decision logic described in Section 3.5 and Chapter 5.

3.4 Application and Dashboard Layer

The application layer is intended to translate AI outputs into information that factory staff can act on directly, without needing to interpret raw model outputs. This includes a clear pass/fail indication for each inspected item, the identified defect type and, where supported, its approximate location on the product; a browsable inspection history; and a dashboard summarizing batch-level pass rates, defect-type frequency, and trends over time. The intention is for this layer to make quality data visible and actionable for factory owners and quality managers, in a way that manual, paper-based inspection records typically do not allow.

3.5 Human-in-the-Loop Decision Support

Karigor AI is deliberately designed around a human-in-the-loop principle rather than fully autonomous decision-making. When the AI model's confidence in a prediction falls below a defined threshold, the item would be flagged for review by a human inspector rather than being automatically passed or failed. This is important for two reasons. First, computer-vision models

can be uncertain or incorrect, particularly for defect types that are underrepresented in training data, for unusual lighting conditions, or for product variations not seen during training. Second, the cost of different error types is asymmetric: an AI system incorrectly accepting a defective product (a false negative) can have more serious downstream consequences — customer complaints, returns, reputational damage — than an unnecessary manual review of an acceptable item (a false positive). Retaining human oversight for uncertain or high-stakes cases is therefore treated as a core design principle rather than an optional feature.

3.6 Key Features

The following features summarize the proposed capabilities of Karigor AI. These are presented as intended design features of the concept and not as capabilities of a deployed or tested product.

- AI-assisted visual inspection of visually observable product defects.
- Pass/fail decision support based on model confidence and predefined thresholds.
- Defect classification where sufficient labelled training data are available for a given product category.
- Optional defect localization (bounding box or segmentation) for supported product and defect types.
- Confidence-based routing of uncertain predictions to human inspectors.
- Digital inspection history and image evidence for each inspected item.
- A quality dashboard summarizing batch and production-level statistics.
- Basic defect-trend analysis to help identify recurring quality issues.
- A user-friendly interface intended for factory staff without specialized AI expertise.
- Compatibility with low-cost, commonly available cameras rather than specialized industrial sensors.
- A choice between local/edge processing and optional cloud synchronization, discussed further in Chapter 5.
- Multilingual and Bangla-friendly interface support, proposed as a future enhancement rather than an initial capability.




Figure 1. Proposed Operational Workflow of Karigor AI

As illustrated in Figure 1, a product arriving at the inspection point is imaged, preprocessed, and analyzed by the AI model. The resulting prediction is then evaluated against a confidence threshold: items classified as acceptable with sufficient confidence proceed directly to the digital quality record; items flagged as a potential defect, or predictions falling below the confidence threshold, are routed to a human inspector for a final decision. All outcomes — whether determined by the AI model alone or confirmed by a human reviewer — are logged to the dashboard, inspection record, and analytics layer, forming the basis for ongoing quality monitoring.

Chapter 4: Target Users, Use Cases and Value Proposition

4.1 Target Customers

Karigor AI is intended primarily for small and medium-sized manufacturers whose products can be meaningfully evaluated through visual inspection. Within a manufacturing organization, the most relevant users are factory owners, quality-control (QC) managers, quality inspectors, production supervisors, and line operators who interact with the inspection point directly. Factory owners and QC managers are the primary decision-makers regarding adoption and would be the main audience for dashboard-level summaries and trend analytics, while inspectors and operators would be the direct day-to-day users of the inspection interface.

4.2 Potential Manufacturing Sectors

Because a single computer-vision model cannot reasonably be expected to inspect every type of product, Karigor AI is not proposed as a universal, one-size-fits-all inspection system. Instead, each deployment is assumed to require product-specific image data, model training or fine-tuning, defined quality criteria, and pilot calibration for the specific defect types relevant to that product. With this understanding, potential manufacturing sectors that may be suitable candidates for a Karigor AI deployment include, as illustrative examples rather than an exhaustive list:

- Garments and textile accessories, where visible stitching, fabric, and surface defects are common quality concerns [4].
- Packaging, including printed cartons, labels, and seals.
- Plastic products, where surface blemishes, discoloration, or moulding defects can occur.
- Leather goods, where surface marks, stitching, and finishing defects are typically inspected visually.
- Handicrafts and small consumer products with visually assessable finish quality.
- Light manufacturing components where surface-level defects are relevant to acceptance criteria.

4.3 Practical Use Cases

Use Case 1: Garment Finishing Inspection

A small garment manufacturer could position a camera at the finishing and packing stage to screen for visible stitching defects, fabric flaws, or asymmetries before items are packed. Items flagged by the system would be set aside for a human inspector to review, while high-confidence acceptable items would proceed directly to packing, with all outcomes logged for later analysis.

Use Case 2: Packaging and Label Quality Check

A packaging workshop could use Karigor AI to check printed cartons or labels for visible misprints, missing elements, or seal defects at the end of a production line, reducing reliance on manual spot-checks and providing a digital record of inspection outcomes for each batch.

Use Case 3: Plastic or Leather Goods Surface Inspection

A small plastics or leather-goods workshop could deploy a single camera station to screen finished items for visible surface defects such as scratches, discoloration, or incomplete finishing, with the system providing an initial pass/fail signal that a supervisor can confirm for flagged or uncertain items before the batch is shipped.

4.4 Value Proposition

For small manufacturers, the intended value of Karigor AI centers on making AI-assisted inspection accessible without requiring the scale, budget, or technical staff typically associated with industrial machine-vision systems. The expected benefits, framed as intended rather than validated outcomes, include:

- More consistent inspection outcomes by reducing reliance on purely subjective human judgment for routine cases.
- Faster initial screening of products, potentially reducing bottlenecks at the inspection stage.
- Digital traceability of inspection outcomes, replacing or supplementing paper-based quality records.

- Earlier identification of recurring or emerging defect patterns through trend analytics.
- A lower-cost, modular entry point into AI-assisted inspection compared with full industrial automation.
- Retention of human expertise and judgment for uncertain or high-stakes inspection decisions.

Table 1. Target Users and Use Cases

User / Role	Primary Interaction	Illustrative Use Case
Factory Owner	Reviews dashboard summaries and adoption decisions	Evaluating overall defect-rate trends across production batches
QC Manager	Monitors quality trends and inspection records	Identifying recurring defect types to inform supplier or process changes
Quality Inspector	Reviews AI-flagged and low-confidence items	Confirming or overriding pass/fail decisions for uncertain cases
Production Supervisor	Oversees the inspection point within the workflow	Coordinating inspection point placement within the production line
Line Operator	Presents products at the inspection point	Placing garments, packaged goods, or components for image capture

Chapter 5: Proposed Technical Architecture and AI Methodology

5.1 Conceptual System Architecture

Karigor AI is conceptualized as a seven-layer system, moving from physical image capture through to optional external integration. Figure 2 presents this proposed architecture, and the sections that follow describe each layer in more detail.



Figure 2. Conceptual System Architecture of Karigor AI

5.2 Image Acquisition

The image acquisition layer consists of a camera positioned at a fixed inspection point, optionally supported by controlled lighting to reduce variability caused by ambient light changes. Consistent, controlled image capture conditions are important for computer-vision reliability, since deep-learning models can be sensitive to lighting and orientation changes not represented in their training data.

5.3 Image Preprocessing

Captured images would be resized, cropped, and normalized before being passed to the AI model. During model development, appropriate data augmentation techniques — such as controlled rotation, brightness adjustment, or cropping — could be applied to training images to improve model robustness to minor variations in imaging conditions, applied only where technically appropriate for the specific defect type and product.

5.4 AI Model Development

The specific computer-vision approach for a given deployment would depend on the nature of the defects being inspected. Image classification could be used where the task is a straightforward acceptable/defective judgment; object-detection architectures such as the YOLO family have shown strong results for localizing and classifying visible surface defects in industrial settings [5], [6], [10]; segmentation approaches could be used where precise defect boundaries matter; and

anomaly-detection techniques may be appropriate where labelled defective examples are scarce relative to acceptable examples.

Given that small manufacturers are unlikely to have access to large, pre-existing labelled defect datasets, transfer learning — fine-tuning a model pretrained on a large general-purpose image dataset using a smaller, product-specific dataset — is proposed as the primary approach to model development. Published research on industrial defect detection has repeatedly shown that transfer learning can achieve meaningful performance even with relatively limited amounts of product-specific training data, which is directly relevant to the data constraints small manufacturers are likely to face [12].

These technical approaches are presented as possible implementation choices to be evaluated during future development, not as models that have already been built, trained, or tested for Karigor AI.

5.5 Decision Engine

The decision engine would apply a confidence threshold to each AI prediction. Predictions above the threshold would proceed to an automatic pass or flag outcome, while predictions below the threshold would be routed to a human inspector, consistent with the human-in-the-loop principle described in Section 3.5. The specific threshold value would need to be calibrated during pilot validation for each product category, balancing the goal of catching genuine defects against the operational cost of excessive manual review.

5.6 Data Storage and Dashboard

Inspection outcomes, associated images, and metadata would be stored in a data layer supporting the dashboard and application layer described in Section 3.4. This data layer underpins inspection history, batch-level summaries, and defect-trend analytics, and would also serve as the basis for periodic model retraining as more validated examples become available.

5.7 Edge vs Cloud Deployment

Karigor AI is conceived to support either local (edge) processing or optional cloud synchronization, depending on a manufacturer's circumstances. Edge inference — running the AI model on a local device at the factory — offers lower latency, continued operation during internet outages, and greater control over where product images are stored, which may be particularly relevant given unstable internet connectivity in some manufacturing locations and general data-privacy considerations. Cloud-based processing, by contrast, may be useful for centralized model updates, remote dashboard access across multiple sites, and reduced local hardware requirements. A hybrid approach — local inference for real-time decisions, with periodic cloud synchronization for dashboards and model updates — is proposed as a reasonable balance for many small-manufacturer deployments, to be refined based on the specific needs of each pilot.

5.8 Proposed Model Evaluation

Any future Karigor AI prototype would need to be evaluated using standard computer-vision performance metrics appropriate to the chosen task. No such evaluation has been conducted, and no numerical results currently exist. Table 2 outlines the metrics that would be relevant to a future evaluation, along with their purpose.

Table 2. Proposed Evaluation Metrics for a Future Karigor AI Prototype

Metric	Purpose	Why It Matters
Precision	Proportion of items flagged as defective that are genuinely defective	High false-positive rates would create unnecessary manual review workload
Recall	Proportion of genuine defects that the model successfully detects	Low recall means defective items may pass through undetected
F1-Score	Balanced measure combining precision and recall	Useful when both false positives and false negatives carry real cost
Accuracy	Overall proportion of correct pass/fail decisions	A general indicator, most meaningful when defect classes are balanced
mAP (mean Average Precision)	Standard detection-quality measure for object-detection models	Relevant if the model localizes defects rather than only classifying images
IoU (Intersection over Union)	Measures overlap between predicted and actual defect location	Relevant for evaluating defect localization or segmentation accuracy
False Accept Rate	Rate at which defective items are incorrectly passed	Directly tied to the most operationally serious type of error
False Reject Rate	Rate at which acceptable items are incorrectly flagged	Affects manual review workload and production efficiency
Inference Time	Time taken for the model to process one image	Determines whether the system can keep pace with production speed
System Latency	End-to-end time from image capture to decision output	Affects usability and integration into the production workflow

5.9 Future Pilot Validation

Beyond controlled technical evaluation, a future pilot would need to compare AI-assisted inspection outcomes against standard human QC decisions on the same set of products, over a defined trial period, at a cooperating manufacturer. This would allow an assessment of not only detection performance but also practical factors such as inspection time per item, the proportion of items requiring manual review, and inspector confidence in the system's flagged decisions. No such pilot has yet been conducted; this is described further as part of the implementation roadmap in Chapter 10.

Chapter 6: Business Model and Market Opportunity

6.1 Business Opportunity

The business opportunity for Karigor AI rests on the gap between the growing technical feasibility of AI-assisted visual inspection and its continued inaccessibility to small and medium-sized manufacturers. Small manufacturers represent a very large share of Bangladesh's industrial base — cottage, micro, small, and medium enterprises together account for the large majority of industrial units and industrial employment in the country [7], [8] — yet most of this segment remains outside the reach of conventional industrial machine-vision providers, whose systems are priced and engineered for large-scale operations. A modular, lower-cost AI-assisted inspection concept, if developed and validated, could address a genuine unmet need within this segment.

6.2 Customer Segments

The primary customer segment is small and growing manufacturers with visually inspectable products, as described in Chapter 4. A secondary segment could include contract manufacturers producing on behalf of larger brands, where consistent, documented quality control may be a customer requirement rather than only an internal operational goal.

6.3 Revenue Model

A modular, recurring-revenue business model is proposed, intended to lower the upfront barrier to adoption relative to full industrial automation while providing a sustainable basis for ongoing model maintenance and support:

- Setup and installation fee — covering camera and lighting setup, initial calibration, and onboarding for a specific product line.
- Recurring software subscription — covering dashboard access, inspection software, and periodic model maintenance.
- Custom model training — for manufacturers requiring inspection of an additional product category or defect type.
- Additional production-line coverage — extending the system to further inspection points within the same facility.
- Technical support and maintenance — ongoing assistance with hardware or software issues.
- Future analytics or API/enterprise integration modules — proposed as longer-term revenue opportunities rather than initial offerings.

6.4 Proposed Service Packages

To illustrate how a modular subscription structure might be organized, three indicative service tiers are outlined below — Starter, Growth, and Professional — differentiated primarily by the number of inspection points, defect categories covered, and level of support included, rather than by specific guaranteed pricing.

Illustrative Business Assumptions for Future Validation — the figures and tiers below are indicative only, intended to demonstrate how a future business model might operate, and are not actual market prices.

Table 3. Proposed Business Revenue Channels

Revenue Channel	Description	Nature of Revenue
Setup / Installation Fee	Camera and lighting setup, calibration, and onboarding for one inspection point	One-time
Starter Subscription (illustrative)	Single product category, basic dashboard, standard support	Recurring
Growth Subscription (illustrative)	Multiple defect types, trend analytics, priority support	Recurring
Professional Subscription (illustrative)	Multiple inspection points, advanced analytics, dedicated support	Recurring
Custom Model Training	Training or fine-tuning for a new product category or defect type	Project-based
Additional Line Integration	Extending coverage to further inspection points	One-time + recurring
Technical Support and Maintenance	Ongoing hardware/software assistance beyond the base subscription	Recurring or on-demand

6.5 Market Opportunity

Small and medium enterprises constitute the overwhelming majority of Bangladesh's industrial base. According to Bangladesh Bank research, the industrial sector's contribution to national GDP

has grown steadily in recent years, with SMEs forming a substantial share of that industrial output [7]. More recent reporting citing SME Foundation data indicates that the cottage, micro, small, and medium enterprise segment accounts for the large majority of the country's economic units and a significant share of national employment and GDP [8]. This scale suggests a large potential addressable base for affordable quality-management tools, even though only a subset of these enterprises are manufacturers with visually inspectable products suited to a concept such as Karigor AI.

At the same time, research on Industry 4.0 adoption among Bangladeshi manufacturers consistently identifies cost, technical-skill gaps, and infrastructure limitations as the dominant barriers to adopting advanced production and quality technologies [2], [3]. This indicates that market opportunity for a concept like Karigor AI depends not only on demonstrating technical capability, but specifically on addressing affordability and ease of adoption — the central design goals of this proposal. No specific market-sizing estimate (such as a total addressable market figure) is presented in this report, since a reliable, verifiable figure specific to AI-assisted visual inspection adoption among Bangladeshi SMEs could not be identified in the available literature; the opportunity is therefore described qualitatively rather than quantified.

6.6 Business Model Canvas

Figure 3 summarizes the proposed Karigor AI business model using the standard nine-block Business Model Canvas structure.




Figure 3. Proposed Business Model Canvas for Karigor AI

Chapter 7: Competitive Positioning and Feasibility

7.1 Existing Quality Inspection Alternatives

Small manufacturers considering quality-inspection improvements generally face three broad alternatives: continuing with manual visual inspection, adopting a high-end industrial machine-vision system, or engaging a general AI consulting firm to build a custom solution. Each of these alternatives has distinct advantages and limitations when evaluated against the constraints of a small manufacturer, as summarized in Table 4.

Table 4. Comparison with Existing Quality Inspection Approaches

Alternative	Advantages	Limitations for Small Manufacturers
Manual Inspection	Flexible, familiar, no specialized equipment required	Subjective, fatigue-prone, weak digital traceability [4]
High-End Industrial Machine Vision	High technical capability and reliability at scale	Expensive, complex integration, designed for large factories [3]
Generic AI Consulting / Custom Build	Fully customizable to specific needs	May require large project budgets and extended development timelines
Karigor AI Concept	SME-focused, modular, human-assisted, affordable entry point	Requires product-specific training data and pilot calibration for each deployment

7.2 Karigor AI Business Differentiation

Karigor AI does not claim to introduce a new computer-vision technique; the underlying methods proposed — transfer learning, CNN-based classification, and object-detection architectures such as YOLO — are established approaches already used across industrial defect-detection research [5], [6], [12]. Rather, the differentiation lies in business and deployment design: assembling these established technical approaches into a package specifically scoped, priced, and simplified for small manufacturers. This differentiation can be summarized as:

“Affordable, modular, human-assisted AI quality inspection designed around the constraints of small manufacturers.”

- Start-small deployment beginning with a single product category and inspection point rather than a full-factory rollout.
- Compatibility with low-cost, commonly available cameras rather than specialized industrial sensors.
- Mandatory human review of uncertain predictions, rather than fully autonomous decision-making.
- Product-specific customization rather than a one-size-fits-all model.
- A choice of local edge processing or cloud synchronization, discussed in Section 5.7.
- Simple, accessible quality analytics designed for non-technical factory staff.
- A deliberate focus on SMEs rather than large, already-automated factories.

7.3 Technical Feasibility

The core technical building blocks proposed for Karigor AI — transfer learning on pretrained CNN or YOLO-family architectures, commodity cameras, and low-cost edge-computing devices — are well established in published computer-vision research and are not experimental or speculative [5], [6], [12]. This suggests reasonable technical feasibility for a future prototype,

though feasibility for any specific product category would still depend on the availability of adequate training images and the visual distinctiveness of the relevant defects.

7.4 Operational Feasibility

The proposed workflow is intended to be simple enough for factory staff to operate with limited training: placing products at a fixed inspection point and reviewing flagged items when prompted. This is a considerably lower operational burden than integrating a full industrial machine-vision system, though it would still require initial staff training and a period of adjustment as inspectors learn to work alongside the AI-assisted workflow.

7.5 Economic Feasibility

The modular adoption path — beginning with a single inspection point and a limited subscription tier — is intended to reduce the initial investment required compared with a full industrial machine-vision installation. However, no return-on-investment figures or cost-recovery timelines are presented in this report, as these would depend on manufacturer-specific factors such as current defect rates, rework costs, and production volume that cannot be reliably estimated without pilot data.

7.6 Data Feasibility

Data availability is likely to be the most significant practical constraint on feasibility. Each product category and defect type would require a representative set of labelled images covering both acceptable and defective examples under realistic imaging conditions. Assembling this dataset is expected to be one of the more time-consuming aspects of any future prototype development, though transfer learning is expected to reduce the volume of data required relative to training a model from scratch [12].

7.7 Scalability

Scaling Karigor AI to additional product categories is not expected to be automatic: each new product or defect type would likely require its own data collection, model adaptation or retraining, and calibration process. Scalability is therefore better understood as a repeatable process for onboarding new product categories, rather than a single model that scales to arbitrary products without additional work.

Chapter 8: SWOT, Risks and Responsible AI

8.1 SWOT Analysis

Table 5 summarizes the strengths, weaknesses, opportunities, and threats relevant to the Karigor AI concept.

Table 5. SWOT Analysis

Category	Key Points
Strengths	Affordable, modular concept; SME-focused design; digital traceability; human-assisted design reduces risk of undetected critical errors
Weaknesses	Requires product-specific training data; needs per-product customization; still an early-stage, unvalidated concept; performance sensitive to imaging conditions
Opportunities	Growing Industry 4.0 and manufacturing digitization interest; large SME base in Bangladesh; potential partnerships with SME associations and hardware vendors; developing local AI ecosystem

Category	Key Points
Threats	Established machine-vision vendors moving downmarket; rapid technology change; data-privacy concerns among manufacturers; potential staff resistance; risk of unreliable predictions from a poorly trained model

8.2 Technical and Business Risks

A number of technical and business risks would need to be actively managed during any future development of Karigor AI, as summarized in Table 6.

8.3 Risk Mitigation

Table 6. Risk and Mitigation Matrix

Risk	Impact	Mitigation
Insufficient defect data	Poor model performance and unreliable predictions	Structured data-collection plan; transfer learning to reduce data requirements [12]
Class imbalance (few defect examples vs. many acceptable items)	Model biased toward predicting “acceptable”	Targeted collection of defect examples; appropriate augmentation and evaluation metrics
Changing lighting conditions	Inconsistent image quality and reduced accuracy	Controlled lighting setup; augmentation during training for lighting variation
Low camera quality	Missed subtle defects	Minimum camera specification guidelines for deployment
Product variation between batches	Model performance may degrade over time	Periodic retraining using validated new examples
False positives (acceptable items flagged)	Unnecessary manual review workload	Threshold tuning during pilot validation
False negatives (defects accepted)	Defective products may reach customers — the most serious risk category	Conservative confidence thresholds; mandatory human review for uncertain cases
Model drift over time	Gradual decline in prediction reliability	Scheduled monitoring and periodic retraining
Hardware failure	Inspection point becomes unavailable	Fallback to manual inspection procedures; basic maintenance plan
Staff resistance to AI-assisted workflow	Low adoption or workaround behavior	Training, transparent communication of AI role, and human-in-the-loop design
Cybersecurity vulnerabilities	Unauthorized access to inspection data or systems	Access controls, secure data handling practices, and regular software updates
Data privacy concerns	Manufacturer or worker discomfort with image capture	Minimizing collection of unnecessary imagery; clear data-ownership policies

Among these risks, false negatives — cases where a defective item is incorrectly accepted by the AI system — are treated as particularly serious, since an undetected defect can reach the customer with consequences more severe than an unnecessary manual review. This is a central reason why Karigor AI's design retains mandatory human review for low-confidence predictions rather than pursuing fully autonomous decision-making.

8.4 Ethical Considerations

Karigor AI is positioned as a decision-support system that assists human quality-control judgment, not as an infallible autonomous authority. This distinction has direct ethical

implications: the system's outputs should be presented to users as recommendations with an associated confidence level, not as unquestionable final verdicts, and human inspectors should retain the ability to override AI decisions, particularly for flagged or uncertain items.

8.5 Privacy and Data Ownership

Because Karigor AI involves capturing images at a production line, careful attention to data privacy is required. Image capture should be limited to product images rather than unnecessarily including workers or unrelated areas of the factory floor. Manufacturers should retain clear ownership of their inspection data, with transparent policies on how images and inspection records are stored, used for model improvement, and, where applicable, secured against unauthorized access.

8.6 Responsible AI and Human Oversight

Consistent with the principles above, responsible use of Karigor AI would require transparency about the AI model's role and limitations, audit logs of inspection decisions for accountability, secure access controls, clearly defined and periodically reviewed confidence thresholds, and a defined process for retraining the model as new validated data become available. Human oversight is treated as a permanent design feature rather than a temporary limitation to be removed as the technology matures.

Chapter 9: Sustainability and Potential Impact

If successfully developed and validated, Karigor AI could plausibly contribute to more sustainable manufacturing practices among small producers, primarily through improved consistency in quality control. Earlier identification of defective products, before further processing or shipment, may help reduce material waste and rework, though the extent of any such reduction has not been measured and should not be assumed without pilot evidence.

The concept may also contribute to improved production efficiency by reducing the time spent on undirected manual re-inspection and by providing quality managers with data to identify and address recurring defect patterns at their source, potentially supporting more responsible and resource-efficient manufacturing practices over time.

These potential contributions can be loosely associated with several United Nations Sustainable Development Goals (SDGs) [11], understood as directional relevance rather than measured impact:

- SDG 8 (Decent Work and Economic Growth) — by supporting the competitiveness and productivity of small manufacturing enterprises.
- SDG 9 (Industry, Innovation and Infrastructure) — by extending access to AI-assisted manufacturing technology to a segment often excluded from such innovation.
- SDG 12 (Responsible Consumption and Production) — by potentially reducing material waste associated with undetected defects and rework.

These associations are presented as plausible, intended directions of impact rather than demonstrated outcomes. Any specific claims about waste reduction, efficiency gains, or environmental benefit would require empirical validation through pilot testing, which has not yet been conducted.

Chapter 10: Implementation Roadmap and MVP

10.1 Development Roadmap

Figure 4 outlines a proposed seven-phase roadmap for developing Karigor AI from its current conceptual stage toward a validated, deployable minimum viable product. No phase beyond the current conceptual stage has yet been carried out, and no dates are assigned, since a realistic timeline would depend on factors such as access to a cooperating manufacturer, data-collection speed, and available development resources.




Figure 4. Proposed Implementation Roadmap for Karigor AI

Table 7. Proposed Implementation Phases

Phase	Description	Key Output
1. Concept Development	Define the use case, target customer problem, and system scope	Validated concept and problem definition (this report)
2. Use-Case Selection	Select one specific product category and a limited set of defect types for initial focus	Defined initial product scope
3. Dataset Development	Collect and annotate representative images of acceptable and defective products	Labelled image dataset
4. MVP Development	Build a basic camera setup, AI model, and dashboard for the selected use case	Working minimum viable product
5. Controlled Validation	Evaluate the MVP using the metrics outlined in Table 2 under controlled conditions	Documented technical performance results
6. SME Pilot	Test the MVP within a real manufacturing environment alongside existing manual inspection	Pilot outcomes and user feedback
7. Commercialization	Refine the product based on pilot findings and introduce the subscription-based business model	Market-ready offering

10.2 Minimum Viable Product

Rather than attempting to build a system capable of inspecting many product types at once, the proposed MVP deliberately narrows scope to a single, controlled use case. This is intended to make initial development realistic and to allow meaningful evaluation before considering any expansion. The proposed MVP would include:

- One specific product category, selected in coordination with a cooperating manufacturer.
- One or a small number of clearly defined defect types relevant to that product.
- A fixed camera position with controlled lighting.
- A basic AI model trained via transfer learning on the collected product-specific dataset.
- Pass/fail output with an associated confidence score.
- Image logging for every inspected item.
- A simple dashboard showing inspection history and basic summary statistics.
- A human-override mechanism allowing an inspector to confirm or reverse any AI decision.

10.3 Data Collection Plan

Data collection for the MVP would involve systematically photographing both acceptable and defective examples of the selected product under the controlled imaging conditions defined for the inspection point, with defective examples further organized by defect type where feasible. This dataset would then be split into training, validation, and test subsets following standard machine-learning practice.

10.4 Prototype Development

Prototype development would involve implementing the image-capture pipeline, training an initial model via transfer learning, building the confidence-based decision logic, and constructing a basic dashboard interface, following the architecture described in Chapter 5.

10.5 Controlled Testing

Before any pilot deployment, the MVP would be evaluated on a held-out test dataset using the metrics defined in Table 2, allowing technical performance to be assessed under controlled conditions before introducing the system into a live production environment.

10.6 SME Pilot

Following successful controlled testing, a small-scale pilot at a cooperating manufacturer would compare AI-assisted inspection against the manufacturer's existing manual process over a defined trial period, as outlined further in the pilot validation discussion in Section 5.9.

10.7 Commercialization Path

Only after controlled testing and pilot validation demonstrate acceptable performance would commercialization — including the subscription-based business model outlined in Chapter 6 — be pursued. This sequencing is intended to ensure that any future commercial offering is grounded in demonstrated, validated performance rather than untested assumptions.

Chapter 11: Limitations and Future Development

11.1 Current Limitations

Karigor AI, as presented in this report, remains at a conceptual stage. The following limitations should be clearly understood by any reader of this report:

- Karigor AI has not been implemented as a working prototype and no software or hardware system currently exists.
- No real-world performance results, accuracy figures, or evaluation metrics have been produced, since no model has been trained or tested.
- No pilot has been conducted with any manufacturer, and no customer relationships currently exist.
- Each future product deployment would require its own product-specific dataset, which has not yet been collected for any product category.
- Model performance, if a prototype is eventually built, would likely be sensitive to imaging conditions such as lighting and camera positioning.
- Different products and defect types would require separate model adaptation, retraining, and calibration rather than a single universal model.
- Any AI model developed is expected to make occasional errors, reinforcing the need for continued human verification of uncertain cases.
- The business model and illustrative pricing presented in Chapter 6 are assumptions for future validation and have not been tested with real customers.

11.2 Future Technical Enhancements

Beyond the initial MVP, several technical enhancements could reasonably be considered in later development phases:

- Additional defect categories and product types beyond the initial MVP scope.
- A multilingual, Bangla-friendly dashboard interface.
- Mobile monitoring access for factory owners and managers.
- Dedicated edge-AI hardware for improved on-site inference performance.
- Integration with IoT sensors for broader production-line monitoring.
- Root-cause defect analytics linking defect patterns to specific production stages.
- Automatic generation of periodic quality reports.
- Continuous model monitoring with alerts for potential performance drift.

11.3 Future Business Expansion

If the concept is validated through prototype development and pilot testing, potential business-expansion directions could include:

- Expansion to additional manufacturing sectors beyond the initial pilot product category.
- API integration with existing ERP or MES systems used by larger SME customers.
- Partnerships with SME associations, chambers of commerce, and hardware suppliers to support distribution and onboarding.
- Development of enterprise-tier offerings for manufacturers with multiple production lines or facilities.

These future directions are presented as plausible expansion paths contingent on successful validation of the core concept, rather than as committed elements of the current proposal. Karigor

AI is not intended to become an unrealistically broad, all-purpose manufacturing AI platform; expansion, if pursued, should remain grounded in demonstrated performance for each new use case.

Chapter 12: Conclusion

Small and medium-sized manufacturers occupy a difficult position with respect to quality control: they face the same fundamental challenges of manual visual inspection — fatigue, subjectivity, and weak digital traceability — as larger manufacturers, yet they are largely excluded from the automated inspection technologies that could help address these challenges, due to cost, technical complexity, and integration requirements designed around large-scale factory operations. This gap between technical feasibility and practical accessibility forms the central problem this report has examined.

Karigor AI has been proposed as a conceptual response to this gap: an affordable, modular, AI-assisted visual quality-inspection concept designed specifically around the constraints of small manufacturers. The proposed system combines a computer-vision model — built using established techniques such as transfer learning and, where appropriate, object-detection architectures like YOLO — with a confidence-based decision engine that routes uncertain predictions to human inspectors rather than pursuing fully autonomous decision-making. This human-in-the-loop design reflects a deliberate judgment that, particularly given the serious consequences of an undetected defect, Karigor AI should function as a tool that augments human quality-control judgment rather than one that attempts to replace it.

The report has outlined a proposed technical architecture spanning image acquisition, preprocessing, AI inference, decision logic, data storage, and dashboard reporting; a development methodology grounded in transfer learning and standard computer-vision evaluation practice; a modular subscription-based business model; and a realistic, phased roadmap moving from the current conceptual stage through dataset development, MVP construction, controlled testing, and SME pilot validation before any commercialization is pursued. Throughout, the report has aimed to distinguish clearly between what is technically plausible based on established research and what has actually been built or validated for Karigor AI — which, at present, is nothing beyond this conceptual proposal.

Karigor AI's potential business and social value — supporting more consistent quality control, reducing reliance on purely manual inspection, and extending AI-assisted tools to a manufacturing segment that has historically been excluded from them — remains, at this stage, a proposed direction rather than a demonstrated outcome. Its ultimate feasibility depends on future work that this report does not claim to have completed: assembling real product-specific datasets, building and testing an MVP, and validating that AI-assisted inspection meaningfully improves on existing manual practice within a real small-manufacturing environment. With that future work, Karigor AI represents a realistic and technically grounded direction for making AI-assisted quality inspection accessible to small manufacturers who currently have no practical path to it.

References

- [1] S. Salman, S. Hasanat, R. Rahman, and M. Moon, “Analyzing the key barriers of adopting Industry 4.0 in Bangladesh's ready-made garment industry: an emerging economy example,” *International Journal of Industrial Engineering and Operations Management*, vol. 6, no. 3, pp. 232–255, 2024.
- [2] S. Hossain, S. Hassan, and R. Karim, “Assessment of critical barriers to Industry 4.0 adoption in manufacturing industries of Bangladesh: An ISM-based study,” *Brazilian Journal of Operations & Production Management*, vol. 20, no. 3, p. 1797, 2023.
- [3] H. M. Elhusseiny and J. Crispim, “SMEs, Barriers and Opportunities on adopting Industry 4.0: A Review,” *Procedia Computer Science*, vol. 196, pp. 864–871, 2022, doi: 10.1016/j.procs.2021.12.086.
- [4] M. B. Ramzan, S. M. Qureshi, M. S. Habib, T. Naveed, M. Awais, and H. Jamshaid, “Evaluation of human factors on visual inspection skills in textiles and clothing: A statistical approach,” *Journal of Engineered Fibers and Fabrics*, vol. 17, 2022.
- [5] E. Cumbajin, N. Rodrigues, P. Costa, R. Miragaia, L. Frazão, N. Costa, A. Fernández-Caballero, J. Carneiro, L. H. Buruberri, and A. Pereira, “A systematic review on deep learning with CNNs applied to surface defect detection,” *Journal of Imaging*, vol. 9, no. 10, p. 193, 2023.
- [6] Y. Gao, G. Lv, D. Xiao, X. Han, T. Sun, and Z. Li, “Research on steel surface defect classification method based on deep learning,” *Scientific Reports*, vol. 14, art. 8254, 2024.
- [7] Bangladesh Bank, “Estimating the Contribution of SMEs Output on GDP Growth in Bangladesh — A VECM Approach,” *Bangladesh Bank Research Paper Series (SRWSP 2022-02)*, Bangladesh Bank, Dhaka, Bangladesh, 2022.
- [8] T. Mohiuddin, “Can SMEs power Bangladesh's next growth chapter?,” *The Business Standard*, Jun. 28, 2026. [Online]. Available: <https://www.tbsnews.net/supplement/can-smes-power-bangladeshs-next-growth-chapter-1473921>
- [9] Dhaka Chamber of Commerce & Industry, “Bangladesh AI Business Summit 2026 – Student Expression of Interest,” DCCI, Dhaka, Bangladesh, 2026. [Online]. Available: <https://aisummit.dhakachamber.com/student/>
- [10] Y. Sun et al., “Parts surface defect detection algorithm based on improved YOLOv8s,” *Engineering Reports*, 2026.
- [11] United Nations, “The 17 Sustainable Development Goals,” United Nations Department of Economic and Social Affairs, 2015. [Online]. Available: <https://sdgs.un.org/goals>
- [12] J. Cheng, B. Guo, J. Liu, S. Liu, G. Wu, Y. Sun, and Z. Yu, “TL-SDD: A transfer learning-based method for surface defect detection with few samples,” in *Proc. 2021 7th International Conference on Big Data Computing and Communications (BigCom)*, Deqing, China, 2021, pp. 136–143, doi: 10.1109/BigCom53800.2021.00023.